



MongoDB, Redis를 연계한 트래픽 처리

hwpark@wiken.co.kr

NoSQL

What is NoSQL?

- NoSQL databases are designed to address specific requirements such as scalability, high availability, and flexible data models.
- Basically a large serialized object store
- In general, doesn't support complicated queries
 - Joins generally not available
- Schema is no longer rigid but flexible
 - Recommends denormalization
- Designed to be distributed (cloud-scale) out of the box
- Drops ACID requirements
 - Any database can answer any query
 - Any write query can operate against any database and will “eventually” propagate to other distributed servers

Types of NoSQL Databases

- Key-Value Store
 - A key that refers to a payload
 - MemcacheDB, Redis, Voldemorte

Key	Value
Name	Joe Bloggs
Age	42
Occupation	Stunt Double
Height	175cm
Weight	77kg

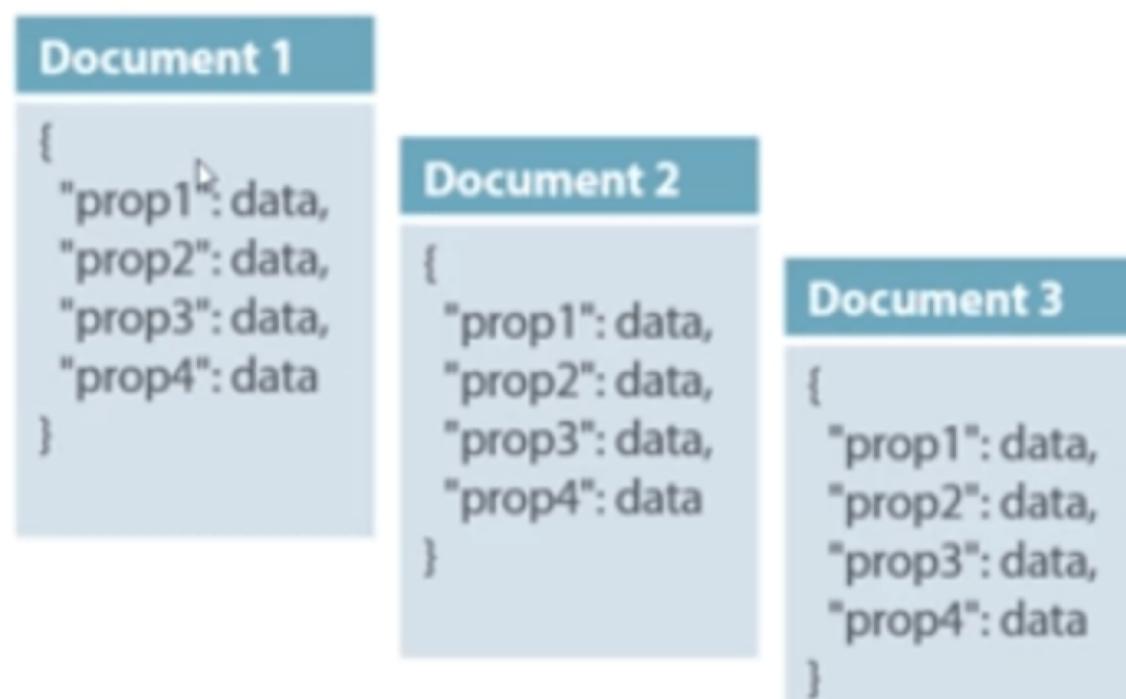
Types of NoSQL Databases

- ColumnStore
 - Column data is saved together, as opposed to row data
 - Super useful for data analytics
 - Hbase, Bigtable, Cassandra

ColumnFamily			
Row Key	Column Name		
	Key	Key	Key
Value	Value	Value	Value
	Column Name		
Key	Key	Key	Key
	Value	Value	Value

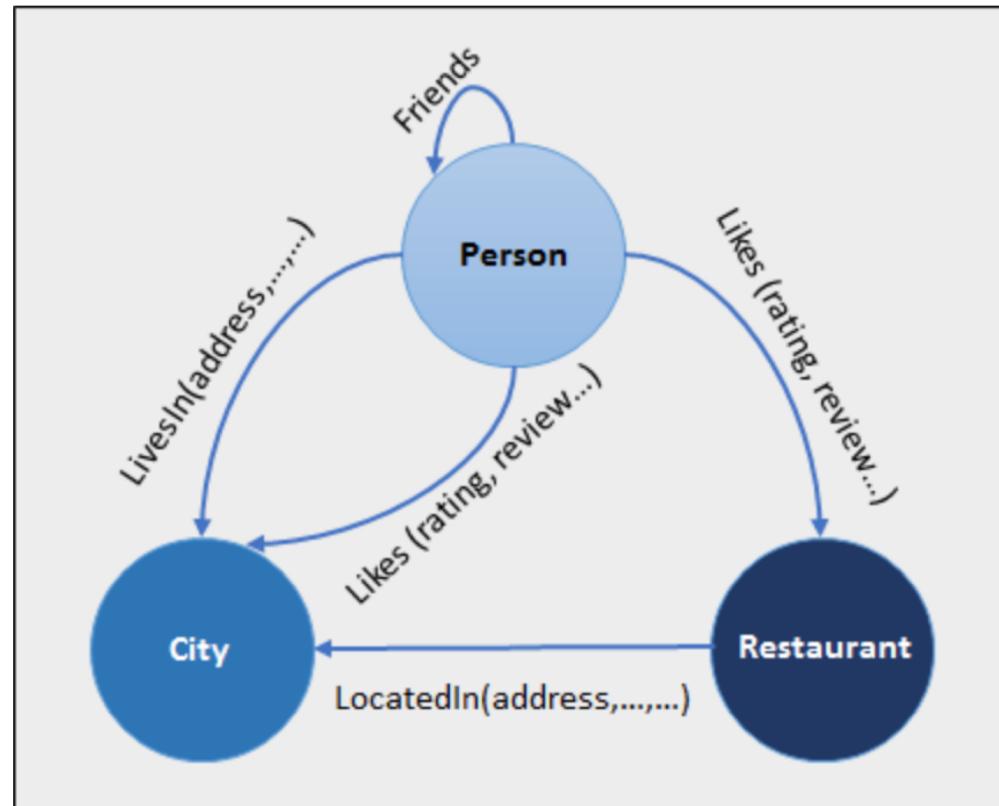
Types of NoSQL Databases

- Document/XML/ObjectStore
 - Key (and possibly other indexes) point at a serialized object
 - DB can operate against values in document
 - MongoDB, CouchDB, Elasticsearch

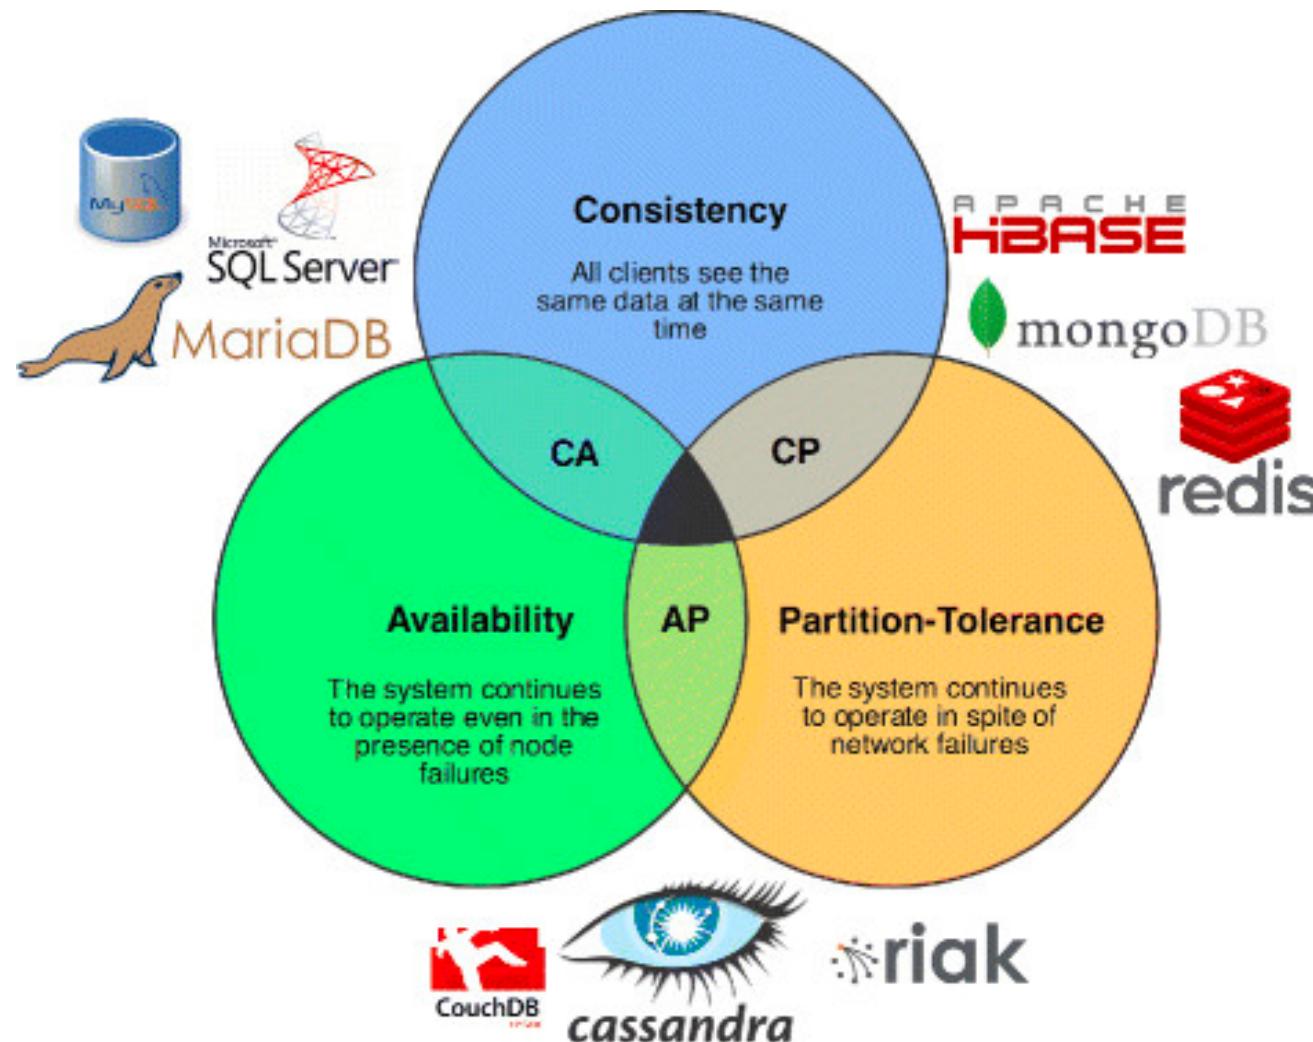


Types of NoSQL Databases

- GraphStore
 - Nodes are stored independently, and the relationship between nodes (edges) are stored with data
 - Neo4j



Types of NoSQL - CAP Theorem Tradeoff



Introduction to MongoDB

- Understanding and Implementing MongoDB
- Welcome to our course "Understanding and Implementing MongoDB & Redis". This module is designed to introduce you to MongoDB, its primary features, use-cases, and how it fits into the modern data management landscape.



The Big Picture

- **High-Level Data Management:**
 - Data management encompasses storage, processing, querying, scalability, and analysis.
 - Consider these aspects when choosing a database system.
- **Traditional Databases vs NoSQL:**
 - Traditional relational databases use a table structure and SQL for data access.
 - They offer strong ACID guarantees but may lack scalability and flexibility.
 - NoSQL databases prioritize scalability, speed, and flexibility over some ACID properties.
 - NoSQL is well-suited for large data volumes and complex data structures.
- **Need for MongoDB:**
 - MongoDB is a NoSQL database with a flexible document-oriented model.
 - It provides scalability through sharding and durability with built-in replication.
 - MongoDB is ideal for content management, user data management, and real-time analytics.

Overview

- **Overview of MongoDB:** We will delve into what MongoDB is, what makes it unique, and why it has gained such popularity in the tech world.
- **MongoDB Use Cases:** We'll then explore a variety of use cases for MongoDB, spanning different industries and applications, to understand where and why it excels.
- **MongoDB Data Model:** We will study MongoDB's unique data model and see how it contributes to MongoDB's flexibility and scalability.
- **Creating and Querying MongoDB Documents:** Finally, we will get hands-on and learn about creating and querying MongoDB documents, the basic operations that form the crux of MongoDB usage

MongoDB Common Terminology

- **_id** – Document's primary key
 - This is a field required in every MongoDB document
 - Unique value in the MongoDB document
 - Automatically created if omitted
- **Collection** – This is a grouping of MongoDB documents
 - Equivalent of a table in other RDBMS such as Oracle or MS SQL
 - A collection exists within a single database
 - Collections don't enforce any sort of structure
- **Cursor** – Pointer to the result set of a query
 - Clients can iterate through a cursor to retrieve results

MongoDB Common Terminology

- **Database** – Container for collections
 - Each database gets its own set of files on the file system
 - A MongoDB server can store multiple databases
- **Document** - Equivalent to a Row in RDBMS
 - A record in a MongoDB collection is basically called a document
 - The document, in turn, will consist of field name and values.
- **Field** - A name-value pair in a document
 - A document has zero or more fields
 - Fields are analogous to columns in relational databases.
- **JSON** – This is known as [JavaScript](#) Object Notation.
 - Human-readable, plain text format for expressing structured data

What is MongoDB?

- **Definition:**
 - MongoDB is an open-source, NoSQL database that uses a document-oriented data model.
 - Unlike relational databases, MongoDB uses BSON (Binary JSON)-like documents and allows for greater flexibility and scalability.



What is MongoDB?

- **Key Features:** Some of the primary features of MongoDB include:
 - **Document-oriented Storage:** Data in MongoDB is stored in flexible, JSON-like documents, meaning fields can vary from document to document and data structure can be changed over time.
 - **Ad hoc Queries:** MongoDB supports a rich set of query operations, including field, range queries, and regular expression searches.
 - **Indexing:** You can index any field in a MongoDB document, significantly improving search performance.
 - **Replication and High Availability:** MongoDB uses native replication for high availability, including automatic failover.
 - **Horizontal Scaling:** MongoDB can be scaled across platforms, allowing for the management of large data sets.
- **Unique Selling Proposition**

Why MongoDB?

- Key Advantages:
 - Schema-less: MongoDB allows collections to hold different documents with varying fields, content, and sizes.
 - Ease of scale-out: Sharding data across multiple servers makes scaling out in MongoDB simple.
 - Rich queries and analytics: MongoDB supports advanced queries, including joins and secondary indexes.
 - Speed: MongoDB uses internal memory for storing the working set, enabling faster data access.

MongoDB Use Cases

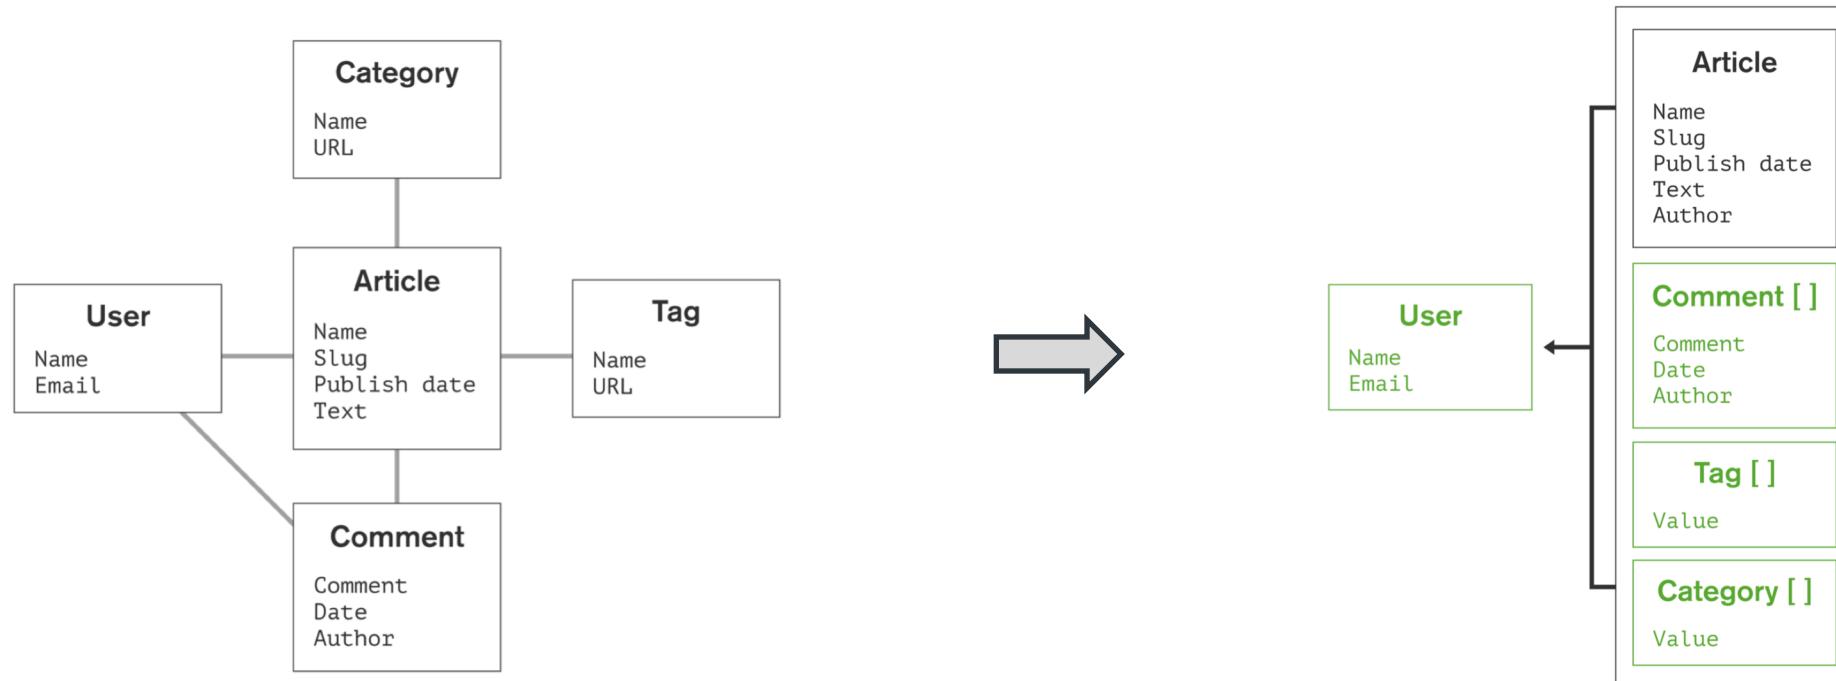
- **Case Studies:** Various organizations across different sectors have successfully integrated MongoDB into their operations:
 - **Healthcare:** MyHealthDirect, a digital care coordination company, leveraged MongoDB to unify patient data and build an event-driven healthcare system.
 - **Finance:** Square, a mobile payment company, used MongoDB to manage its enormous database with high demand for write loads.
 - **Retail:** Urban Outfitters uses MongoDB to keep track of its diverse set of products, customers, and recommendations, all in real time.
- **Success Stories:**
 - **Adobe:** Adobe Experience Manager uses MongoDB to store, manage, and access billions of digital assets, providing a seamless customer experience.
 - **InVision:** InVision built its design prototyping tool using MongoDB due to its ability to handle large and complex data sets and schemas.

MongoDB Data Model

- **Document Model:** MongoDB uses a flexible document model that can accommodate complex hierarchical relationships and an array of data types.
- **BSON Format:** MongoDB stores data in a format known as BSON (Binary JSON). BSON extends the JSON model to provide additional data types, ordered fields, and to be efficient for encoding and decoding within different languages.
- **Collections and Databases:** in MongoDB, a collection is a group of MongoDB Documents and is the equivalent of an RDBMS table. A collection exists within a single database. Collections do not enforce a schema and documents within a collection can have different fields.

MongoDB Data Model – Data as Documents

- MongoDB stores data as documents in a binary representation called BSON (Binary JSON).
- The BSON encoding extends JSON to include additional types
 - int, long, date, and floating point
 - BSON documents contain one or more fields, and each field contains a value of a specific data type, including arrays, binary data and sub-documents.



Creating MongoDB Documents

- **Creating Documents:** MongoDB stores structured data as BSON documents, i.e., data structures composed of field and value pairs
- **InsertOne() and InsertMany() Methods:** MongoDB provides two primary methods for inserting data into a collection: insertOne() and insertMany()

```
db.collection('users').insertOne({  
  name: 'John Doe',  
  age: 28,  
  email: 'johndoe@example.com'  
})
```

```
db.collection('users').insertMany([  
  {  
    name: 'John Doe',  
    age: 28,  
    email: 'johndoe@example.com'  
  }, {  
    name: 'Jane Doe',  
    age: 27,  
    email: 'janedoe@example.com'  
  }])
```

Querying MongoDB Documents

- **Find() Method:** MongoDB uses the find() method to query documents within a collection.

```
db.collection('users').find()  
db.collection('users').find({ age: 28 })
```

- **Query Types:** MongoDB offers a variety of query types, including exact match, range, and logical AND/OR operations.

```
// Exact Match  
db.collection('users').find({ name: 'John Doe' })  
// Range  
db.collection('users').find({ age: { $gt: 20, $lt: 30 } })  
// Logical AND  
db.collection('users').find({ $and: [{ age: { $gt: 20 } }, { age: { $lt: 30 } }] })
```

Querying MongoDB Documents

- **Projections:** Projections in MongoDB are a way to specify which fields you want returned in the output of your `find()` queries.

```
db.collection('users').find({}, { name: 1, email: 1 }) // Only returns the 'name' and 'email' fields
```

Query Example – find()

- Query all employee names with salary greater than 18000 sorted in ascending order

```
db.users.find({salary:{$gt:18000}, {name:1}}).sort({salary:1})
```

Collection Condition Projection Modifier

{salary:25000, ...}
{salary:10000, ...}
{salary:20000, ...}
{salary:2000, ...}
{salary:30000, ...}
{salary:21000, ...}
{salary:5000, ...}
{salary:50000, ...}

→

{salary:25000, ...}
{salary:20000, ...}
{salary:30000, ...}
{salary:21000, ...}
{salary:50000, ...}

→

{salary:20000, ...}
{salary:21000, ...}
{salary:25000, ...}
{salary:30000, ...}
{salary:50000, ...}

Query Example – insert()

- Insert a row entry for new employee Sally

```
db.users.insert({  
    name: "sally",  
    salary: 15000,  
    designation: "MTS",  
    teams: [ "cluster-management" ]  
})
```

Query Example – Update()

- All employees with salary greater than 18000 get a designation of Executive

```
db.users.update(  
    Update Criteria      {salary:{$gt:18000}},  
    Update Action        {$set: {designation: "Manager"}},  
    Update Option        {multi: true}  
)
```

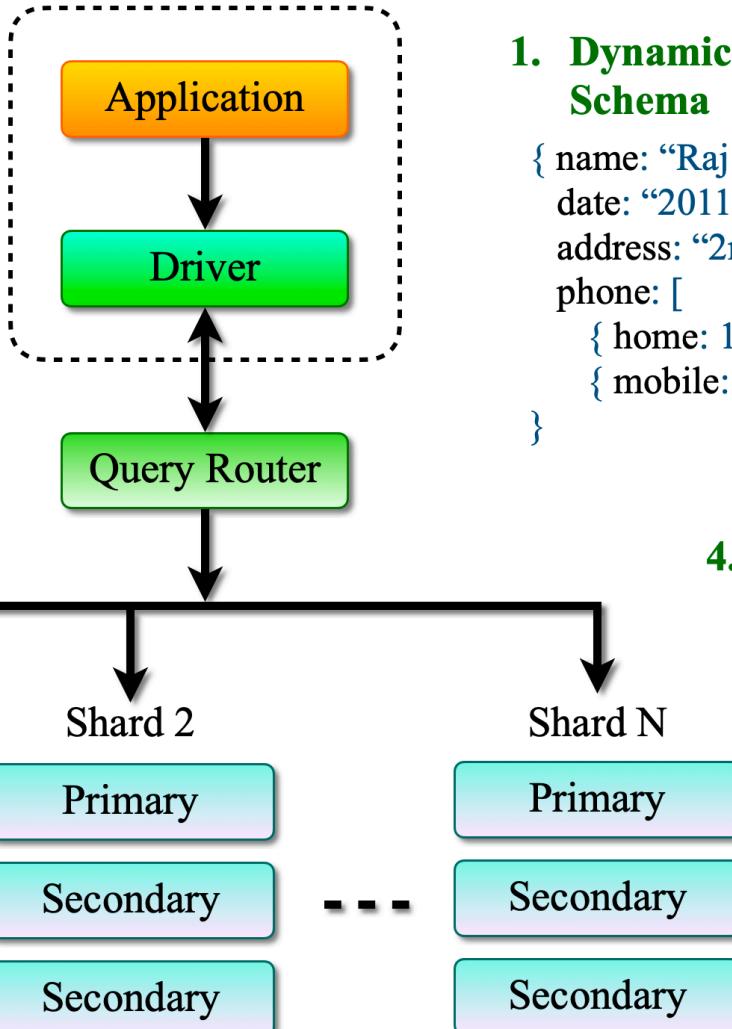
- Multi option allows multiple document update

MongoDB Architecture



2. Native language drivers

```
db.customer.insert({...})  
db.customer.find({  
    name: "Raj Kumar"})
```



3. High availability

- Replica sets

1. Dynamic Document Schema

```
{ name: "Raj Kumar",  
  date: "2011-07-18",  
  address: "2nd Floor.",  
  phone: [  
    { home: 1234567890},  
    { mobile: 1234567890} ]  
}
```

4. High Performance

- Data locality
- Rich Indexes
- RAM

5. Horizontal Scalability

- Sharding

MongoDB Architecture

- **Cluster Architecture:** MongoDB can be run in a clustered environment that offers high availability and horizontal scalability. The cluster architecture includes MongoDB instances running as a part of replica sets or sharded clusters.
- **Replica Set:** A replica set in MongoDB is a group of MongoDB servers that maintain the same data set. Replica sets provide redundancy and high availability, which are achieved by replicating data across multiple servers

```
js                                     Copy code

// Configuring a replica set
mongod --port 27017 --dbpath /data/db --replSet rs0
```

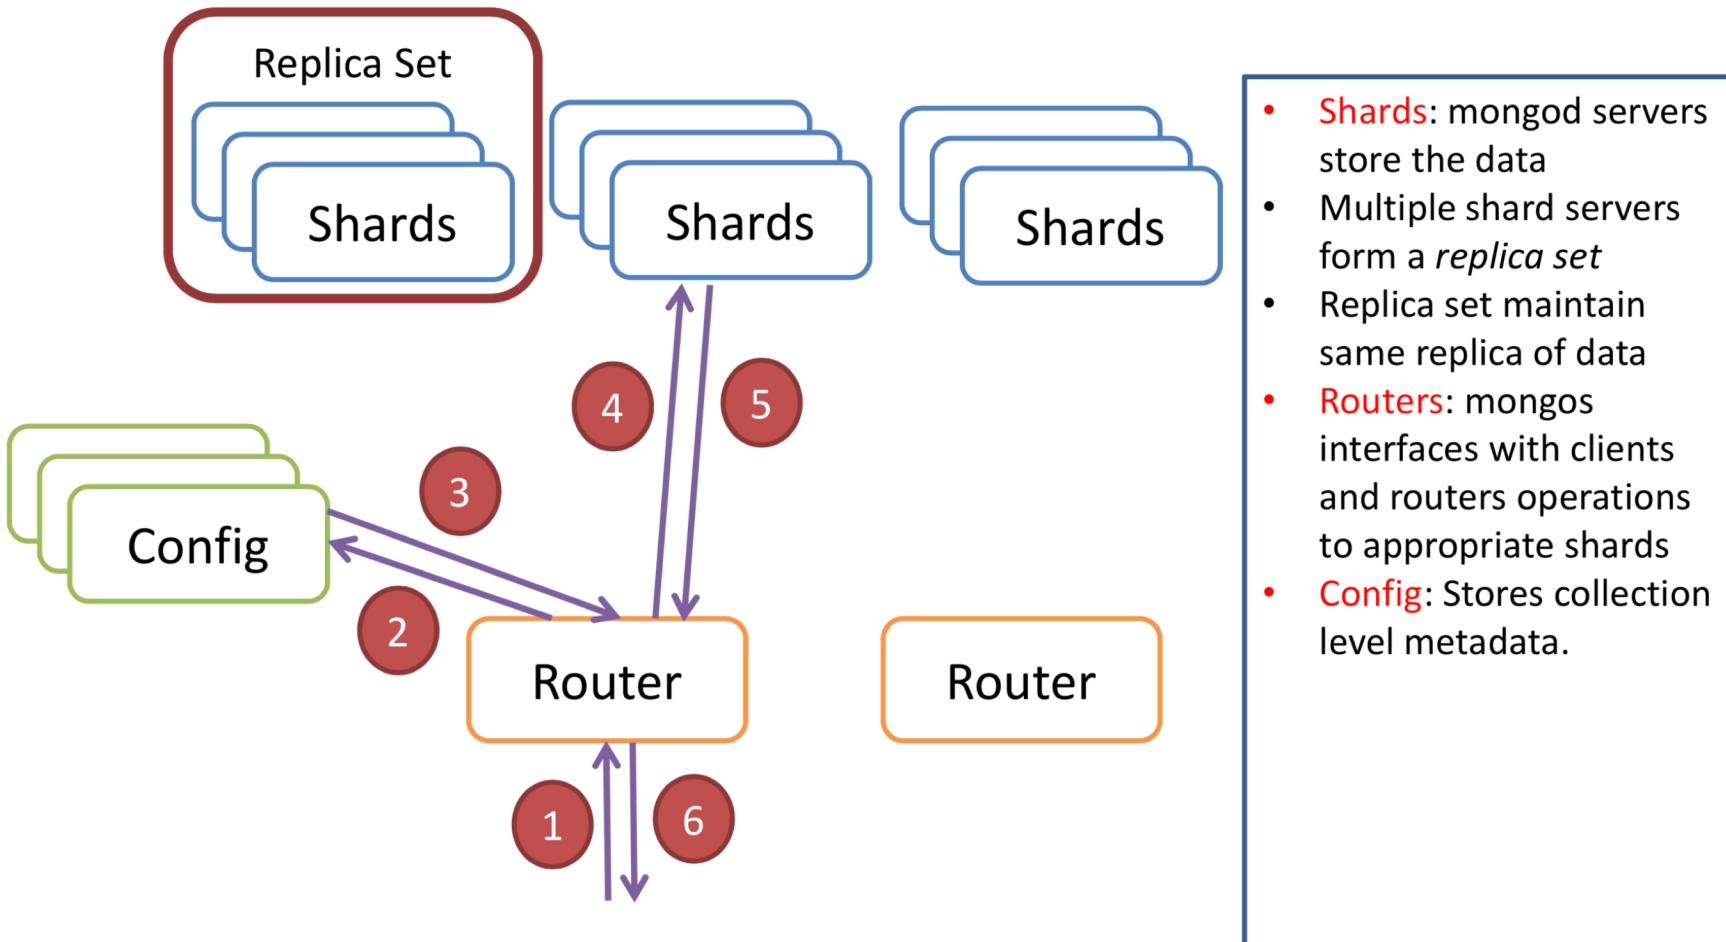
- **Sharding:** Sharding is a method for distributing data across multiple machines.

Replica Set

- A Replica Set in MongoDB is a group of MongoDB instances that host the same data set.
- In a replica set, one node is a primary node that receives all write operations while other nodes, known as secondary nodes, replicate the primary's oplog and apply the operations to their data sets.
- **Components of Replica Set:**
 - The primary node: Processes all write operations.
 - Secondary nodes: Mirror the primary node to provide redundancy and backup.
 - Arbiter node (optional): Participates in elections to help avoid a tie.
- **Advantages of Replica Set:**
 - High data availability
 - Data redundancy
 - Read scalability

Sharding

- Sharding is a database architecture that partitions, or shards, data across multiple servers, providing a method for horizontal scaling.



MongoDB indexing and Aggregation

MongoDB Indexing and Aggregation

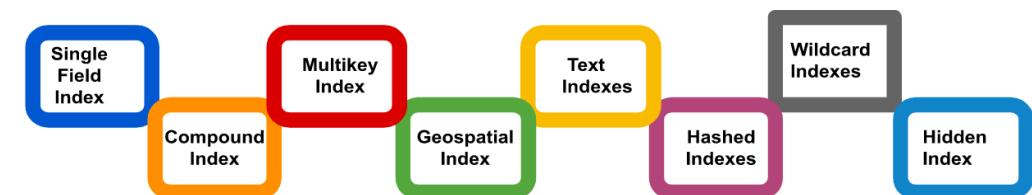
- **Indexing Basics:**
 - Introduction to the concept and role of indexing in MongoDB.
- **Benefits of Indexing:**
 - Dive into how indexing enhances query performance and write speeds.
- **Creating and Using Indexes:**
 - Learn practical skills for creating, using, and managing indexes.
- **Introduction to MongoDB Aggregation Framework:**
 - Understand how MongoDB allows sophisticated data manipulation beyond basic CRUD operations.

Indexing

- **Indexing is like a database 'table of contents'.**
 - Swiftly locates data, eliminating the need for full table scans.
- **Why Indexing in MongoDB?**
 - Crucial for handling vast volumes of data.
 - Significantly improves performance by limiting search space.
- **Types of Indexes in MongoDB:**
 - **Single Field**
 - **Compound Index**
 - **Multikey Index**
 - **Text Index.**
 - **2D Index**

Types of Indexes in MongoDB

- **Single Field Index**
 - Indexing on a single field, including embedded or top-level fields.
 - Enhances efficiency for searches on that field.
- **Compound Index**
 - Indexing on multiple fields.
 - Field order is essential; influences sorting and querying.
- **Multikey Index**
 - Ideal for fields containing arrays of values.
 - Each element of the array is indexed, facilitating specific value searches.
- **Text Index**
 - Suitable for working with text strings.
 - Enables specific text searches within large text blocks
- **Geospatial (2D) Index**
 - Suitable for geospatial queries involving 2D geometries
 - Enhances efficiency of geographical searches.



Single Field Index

- **Single Field Index in MongoDB**
 - Quick access to data based on one field in the document collection.
 - Analogous to a book's index leading to specific page numbers.
- **Creating a Single Field Index**
 - Use `createIndex()` function.
 - Specify the field to be indexed.
 - Example: `db.users.createIndex({name: 1})`.
 - '1' indicates ascending order, '-1' for descending.
- **Use Cases & Performance Benefits**
 - Ideal for queries searching values on a specific field.
 - Reduces the number of documents MongoDB scans, improving performance.
 - Considerations: Indexes require storage space.

Compound Index

- **Compound Indexes in MongoDB**
 - - Index multiple fields together for quicker multi-field search.
 - - Supports queries matching multiple fields, enhancing read operation speeds.
- **Creating a Compound Index**
 - - Use `createIndex()` function.
 - - Specify multiple fields.
 - - Example: `db.users.createIndex({lastName: 1, firstName: 1})`.
 - - '1' indicates ascending order for both fields.
- **Use Cases & Performance Considerations**
 - - Ideal for queries matching on multiple fields, supports sorting of data.
 - - The field order in the index determines the sort order of results.
 - - Considerations: Compound indexes require more storage space.

Multikey Indexes

- **Multikey Indexes in MongoDB**
 - Indexes fields with array of values.
 - Each array element is indexed separately, enabling efficient searches over array contents.
- **Creating a Multikey Index**
 - Similar to creating a single field index.
 - Example: `db.blogPosts.createIndex({tags: 1})`.
 - If 'tags' is an array, MongoDB automatically creates a multikey index.
- **Use Cases & Performance Considerations**
 - Ideal for queries finding documents based on elements within an array field.
 - Considerations: Require more storage due to individual array element indexing. Use judiciously based on application needs.

Text Indexes

- **Text Indexes in MongoDB**
 - Efficiently search for string content within collections.
 - Supports querying of text content.
- **Creating a Text Index**
 - Use `createIndex()` function.
 - Different parameter for text indexing.
 - Example: `db.blogPosts.createIndex({content: "text"})`.
- **Use Cases & Text Search Operations**
 - Ideal for searches for words or phrases within text content.
 - MongoDB provides operators for phrase search, word exclusion, case-sensitive searches, and more.

Geospatial Index

- **2D Indexes in MongoDB**
 - Designed to handle geospatial data, supporting queries on planar geometry.
 - Ideal for data representing points on a map.
- **Creating a 2D Index**
 - Similar to creating other indexes.
 - Example: db.locations.createIndex({coordinates: "2d"}).
 - 'coordinates' field holds an array of [x, y] coordinates.
- **Use Cases & Performance Considerations**
 - Useful for location-based queries, e.g., finding locations within a certain distance from a point.
 - Considerations: Require additional storage. Use only when necessary to support application's requirements.

Benefits of indexing

- **Improved Query Performance:**
 - Reduces the number of documents MongoDB needs to scan.
 - Enhances speed and efficiency of queries.
 - Analogy: Easier to find a book in a well-organized library.
- **Faster Write Operations:**
 - While known for impact on read operations, also improves write operations.
 - Helps MongoDB locate documents faster for updates and inserts.

Creating and Managing Indexes

- **Creating Indexes:**
 - Use `createIndex()` function.
 - Example: `db.users.createIndex({name: 1})`, '1' indicates ascending order.
- **Indexing Strategies:**
 - Understand application's data access patterns.
 - Identify frequently queried fields for indexing.
 - Field order in compound indexes impacts sort order and match conditions.
 - Consider read-write ratio; excessive indexes could slow down heavy-write app.
- **Index Management:**
 - Use `db.collection.getIndexes()` to list all indexes.
 - Remove unnecessary indexes with `db.collection.dropIndex()`.
 - Leverage MongoDB Management Service (MMS) for automatic index suggestions.

Aggregation Framework

- **Concept of Data Aggregation:**
 - Transforming input data into grouped, summarized output.
 - Allows calculations like averages, sums, max, min, etc.
- **Aggregation Pipeline:**
 - Processes data records and returns computed results.
 - Similar to Unix pipeline; output of one operation is input for the next.
 - Can utilize indexes to boost performance.
- **Map-Reduce:**
 - Data processing paradigm for condensing large volumes of data.
- **Single Purpose Aggregation Methods:**
 - Convenience methods like count(), distinct(), and group().
 - Ideal for simple aggregations, but lack flexibility of Pipeline and Map-Reduce.

Aggregation Pipeline

- **Understanding Aggregation Pipeline:**
 - A framework for data transformation and analysis.
 - Processes data records and yields computed results.
 - Works much like a Unix pipeline.
- **Pipeline Operators:**
 - Work on data and modify it.
 - Examples include \$match, \$group, \$sort, and more.
- **Pipeline Execution:**
 - Pipelines run in sequence, with the output of one stage serving as input for the next
 - Allows for efficient data processing within MongoDB.
-

Map-Reduce Function

- **Definition of Map-Reduce:**
 - A technique for processing large data sets.
 - Composed of a 'Map' operation that filters data, and a 'Reduce' operation that performs a summary.
- **Functionality of Map-Reduce:**
 - 'Map' step: Each input document turns into more than one key-value pairs.
 - 'Reduce' step: Key-value pair data is condensed into a single output.
- **Usage Scenarios for Map-Reduce:**
 - Ideal for complex data processing tasks.
 - Use when Aggregation Pipeline doesn't provide required functionality.

Single Purpose Aggregation Methods

- **Count:**
 - Returns the count of documents that match a query.
 - Example usage: db.collection.count(query).
- **Distinct:**
 - Finds the distinct values for a specified field across a single collection.
 - Example usage: db.collection.distinct(field, query).
- **Group:**
 - Groups documents by a specified identifier expression and applies the accumulator expression(s) to each group.
 - Example usage: db.collection.group({ key, reduce, initial}).

Group Method

- **Definition of 'Group':**
 - MongoDB operation that groups documents by a specified identifier.
 - Applies accumulator expressions to each group.
- **Functionality of 'Group':**
 - Allows aggregation tasks like summing field values, averaging, etc.
 - Example usage: db.collection.group({ key, reduce, initial }).
- **Usage Scenario for 'Group':**
 - Ideal for grouping data by certain criteria and performing calculations.
 - For example, finding the total sales by each salesperson.

Summary

- **Review of Key Concepts:**
 - Indexing in MongoDB enhances query performance.
 - We learned about different types of indexes (Single Field, Compound, Multikey, Text, 2D, etc.)
 - We explored MongoDB's Aggregation Framework, including the Aggregation Pipeline, Map-Reduce, and Single Purpose Aggregation Methods.
- **Applying the Knowledge:**
 - We will now apply these concepts in the following lab session.
 - Practical implementation of these techniques will solidify your understanding.

Lab

MongoDB indexing and Aggregation

